





The Transparent Machine

Decoding Feature Importance in Random Forests & Decision Trees

From Scikit-Learn API Calls to Mathematical Architecture


The Interpretability Mandate

High-performing models are useless if their decisions cannot be trusted. When a client asks why a loan was rejected, Feature Importance provides the exact variables that drove the algorithm's decision. It bridges the gap between machine logic and human reasoning.




Redefining the Feature

To understand importance, we must look at the data through the algorithm's eyes.
In the MNIST dataset, an image is not a picture; it is an array of 784 distinct mathematical features.




Raw Image View



1 Image = 1 Row
784 Pixels = 784 Columns (Features)

What the Algorithm Sees


The Random Forest natively detects structural reality. Feature importance isolates the mathematical signal from the noise, identifying the pixels with the highest discriminatory power.



The Scikit-Learn Implementation

Three lines of code yield an array summing to 1.0. But what mathematics generated these numbers?



```
from sklearn.ensemble import RandomForestClassifier  
  
rf = RandomForestClassifier()  
rf.fit(X, y)  
  
importances = rf.feature_importances_
```



The Array Output: Highlighting critical features.


Dissecting the Ensemble

To understand how a Random Forest assigns value to a feature, we must look at its fundamental unit. The architecture of importance begins inside the nodes of a single Decision Tree.




The Currency of Splitting: Impurity

Decision Trees build themselves by finding the feature that best separates mixed data into pure classes. Gini Impurity measures the chaos in a node. A perfect split reduces chaos to zero.




High Impurity
Mixed Data Classes



Low Impurity
Gini = 0.0


Calculating Node Importance

The importance of a single split is calculated by how much it reduces total impurity, weighted by the proportion of samples that pass through it.



The Impurity Funnel in Action

A feature's value is the exact mathematical chaos it removes from the system.




$$(5/5 * 0.48) - (4/5 * 0.375) - (1/5 * 0) = 0.18$$

Importance of this split = 0.18


Tree-Level Normalisation

A single feature may be used multiple times across different branches. The algorithm sums the importance of every node where a feature was used, then normalises the final scores so the entire tree equals 1.0.




Scaling to the Forest

A single tree's perspective is highly dependent on specific training data.
To find the true, stable signal, we must consult the entire ensemble.



The Single Tree

- High Variance
- Prone to Overfitting
- Unstable Signal




The Ensemble


- Stable & Generalised
- Averaged Perspectives
- Reliable Importance

The Ensemble Averaging Engine

The Random Forest aggregates the normalised arrays from every estimator in the ensemble. The final output is simply the arithmetic mean.




The Architecture of Importance



The High-Cardinality Trap

Impurity-based importance has a mathematical bias: it artificially inflates the importance of high-cardinality features (like IDs).

The algorithm easily creates pure splits by fracturing granular categories, creating a misleading interpretation.



Choosing the Right X-Ray

When structural bias threatens interpretability, Scikit-Learn offers `permutation_importance` as a robust alternative to the default Gini calculation.

Gini (Impurity) Importance [Default]	Permutation Importance [Alternative]
In-built during model training	Shuffling data post-training
Extremely Fast	Computationally Expensive
Biased by High Cardinality	Agnostic to Cardinality
Initial rapid diagnostics	Rigorous, final model auditing